

The IMDb logo, consisting of the letters "IMDb" in a bold, black, sans-serif font, set against a yellow rectangular background.

IMDb

Spoiler Detection

Kevin Chiv

Why Movie Spoiler Detection

- Users -



Less Spoilers

- Site Admin -



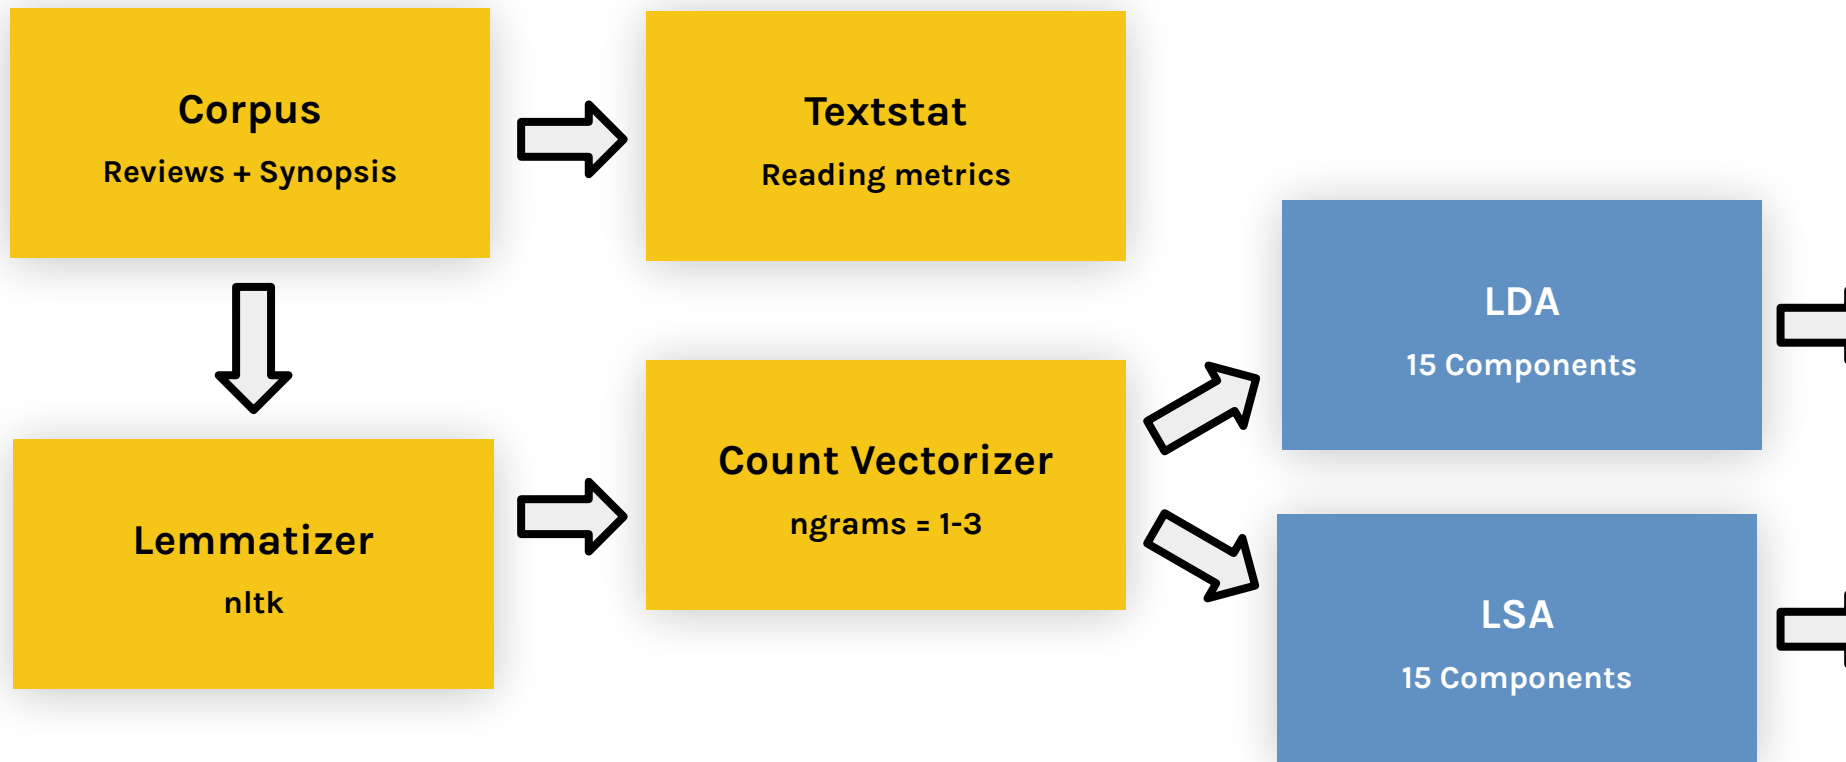
Less to Moderate

- Film Makers -

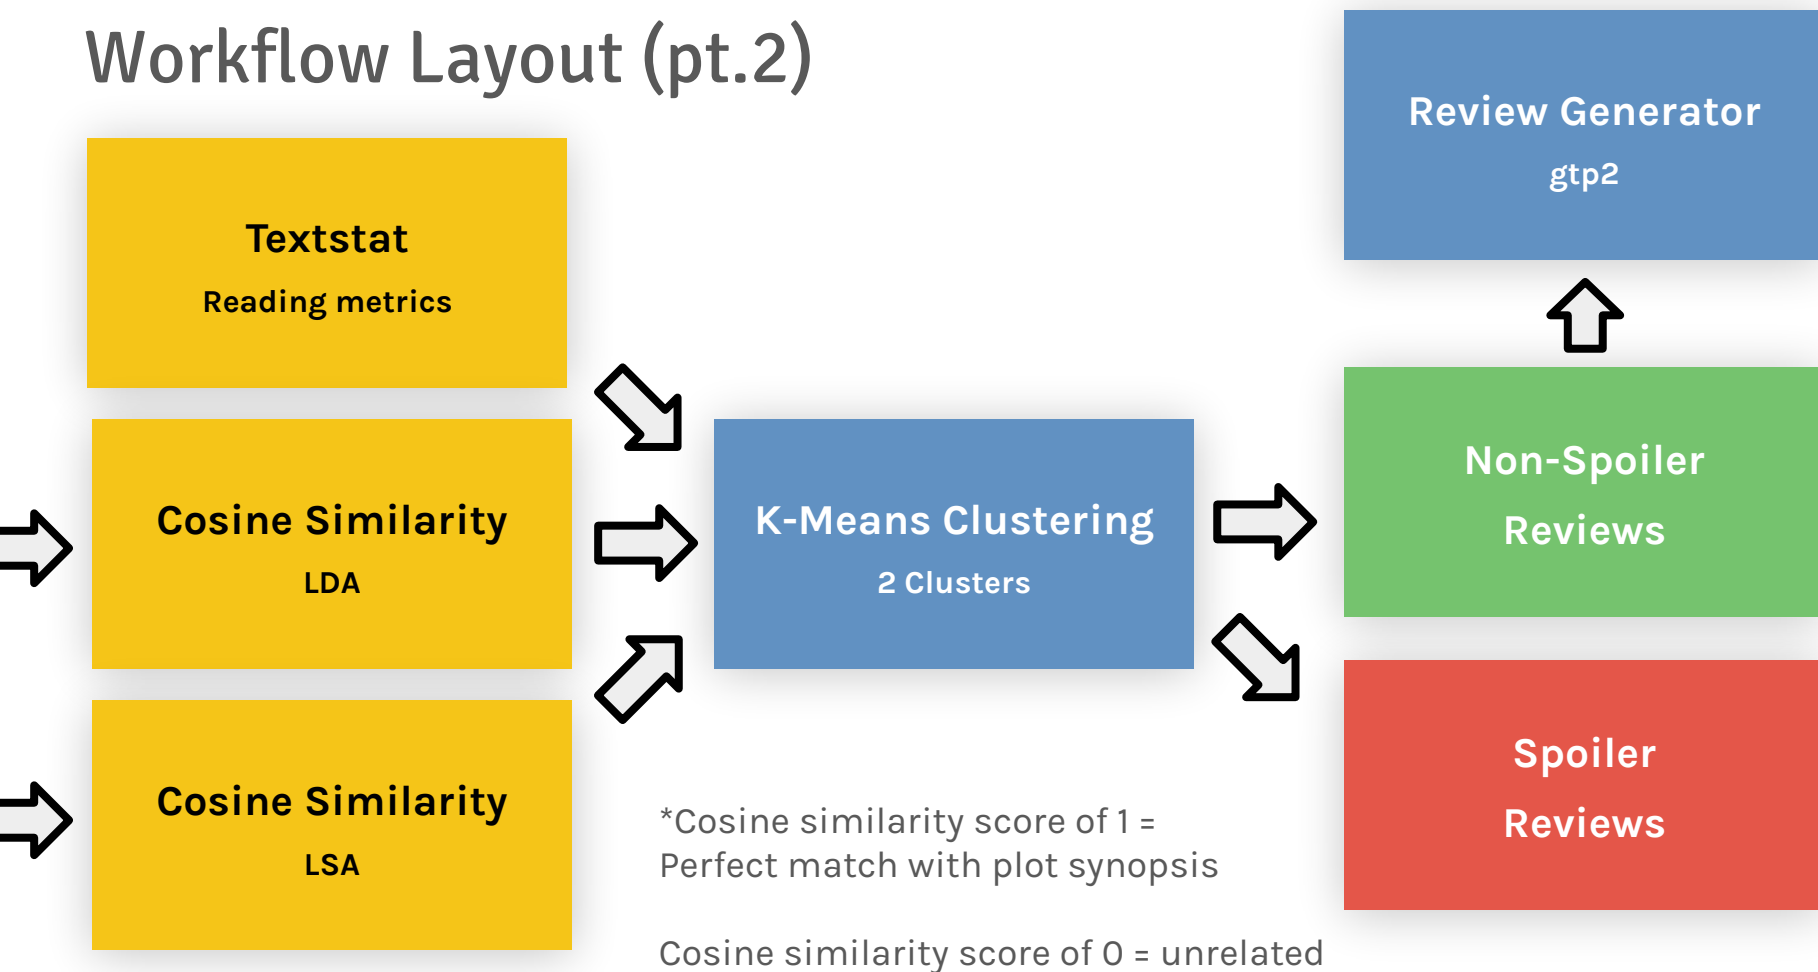


More Revenue

Workflow Layout (pt.1)



Workflow Layout (pt.2)




Let's Look At An Example




Explained Variance

Explained Variance by LSA for Lord of the Rings Reviews



Cumulative
Explained Variance


Cumulative Explained Variance by LSA for LOTR Reviews




Intertopic Distance Map (via multidimensional scaling)



Marginal topic distribution



Top-30 Most Relevant Terms for Topic 12 (7.3% of tokens)



Overall term frequency

Estimated term frequency within the selected topic


1. $\text{saliency}(\text{term} | w) = \text{frequency}(w) * [\sum_t p(t | w) * \log(p(t | w) / p(t))]$ for topics t ; see Chuang et. al (2012)

2. $\text{relevance}(\text{term} | w | \text{topic } t) = \lambda * p(w | t) + (1 - \lambda) * p(w | t) / p(w)$; see Sievert & Shirley (2014)

Movie Review Statistics

	Spoilers	Non-Spoilers
Cosine Similarity (LSA)	0.854	0.610
Cosine Similarity (LDA)	0.804	0.700
Cosine Similarity (NMF)	0.752	0.535
Word Difficulty Score	0.125	0.058
Count	221957	166083

UMAP Projection of Review Statistics



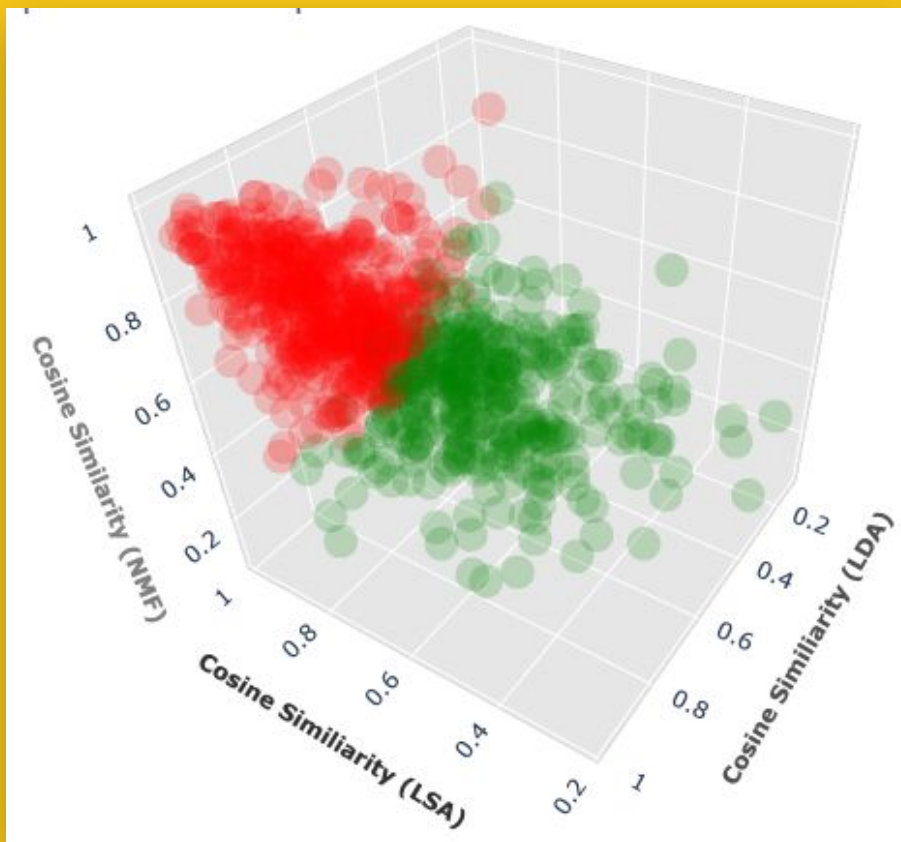
K-Means Clustering

Non-Spoiler Reviews

- Green
- Lower cosine similarity values

Spoiler Reviews

- Red
- About 57% of all reviews



Generated Non-Spoiler Review for LOTR

1000 iterations later ... gpt2 presents* ...

With an an average loss of 1.7 ...

“Just got home from seeing "The Return of the King" at the midnight screening. The movie was great, I had expected a bit more but still felt that there was too much to cover. The action was fantastic, everything from the acting to the special effects. I would recommend this movie to anyone, especially those who haven't read the books...”

**Thanks to Seth Weiland*

Recommendations

- Automatically flag text as spoilers with high cosine similarity against plot synopsis
- Hide spoiler reviews behind a “Click to See Spoiler” button
- Generate spoiler reviews to further improve the model

Arrow - #781 - ██████████ - 6/6/18 - prePINK REVISIONS 1.

ACT ONE

1 EXT. ██████████ - DAY (D1) 1

TRACK across the ██████████ to the unforgiving locale of ██████████

OLIVER'S ██████████ FAMILIAR VOICE speaks to us ██████████

OLIVER (V.O.)

The name ██████████

(beat)

██████████

This recitation ██████████ takes us to --

2 ██████████ SUPERMAX - DAY (D1) 2

Per the end of last episode, ██████████

██████████ The two men separated ██████████ cell. ██████████

██████████

I'm ██████████

OLIVER

I need ██████████

██████████ from his pocket and ██████████ as the Supermax's GUARD (dressed in an ██████████

██████████ storms up --

SUPERMAX GUARD

What the hell's ██████████

OLIVER

There's a problem. ██████████

Before Oliver ██████████ the Guard GRABS his ██████████

The background is a solid yellow color. In the top-right corner, there is a black film strip with white rectangular sprocket holes. The numbers 13, 14, and 15 are visible on the strip. In the bottom-left corner, there is another black film strip with white rectangular sprocket holes. The numbers 15 and 16 are visible on this strip.

~~The End~~
Thank You